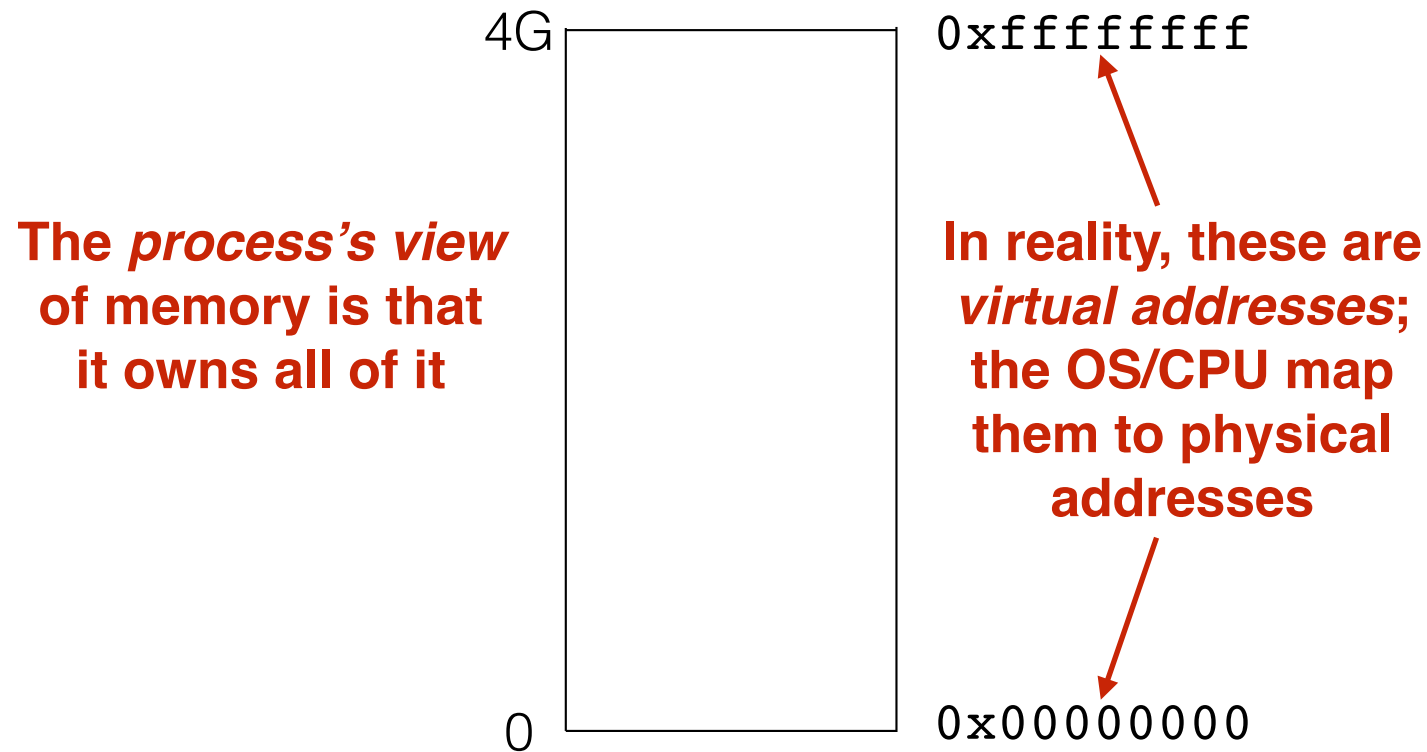


# Memory layout

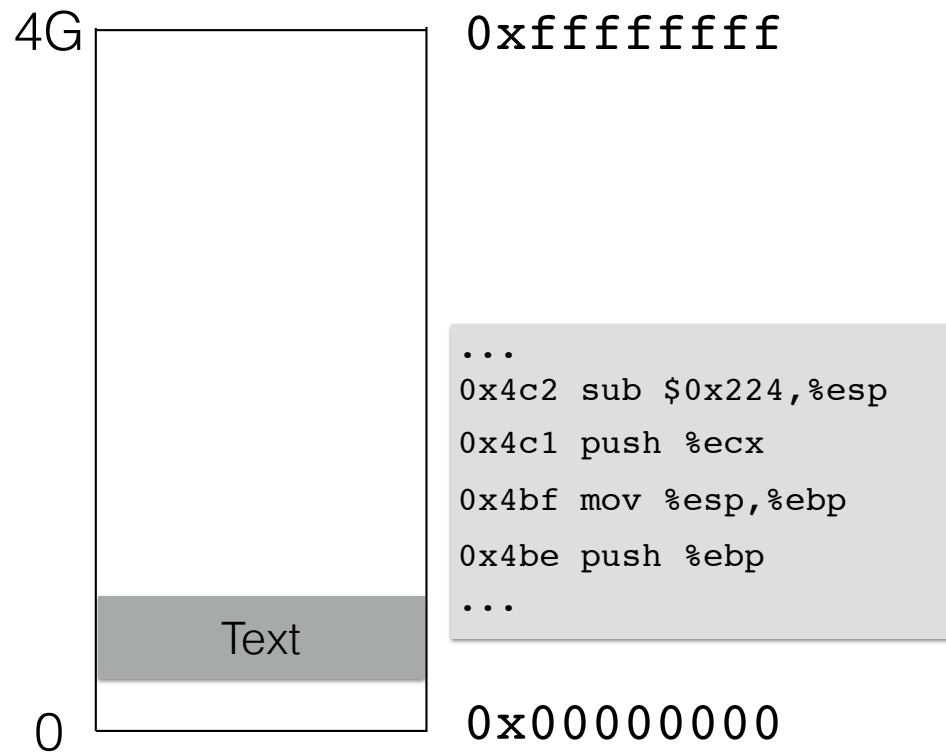
# Memory Layout Refresher

- How is program data laid out in memory?
- What does the stack look like?
- What effect does calling (and returning from) a function have on memory?
- We are focusing on the Linux process model
  - Similar to other operating systems

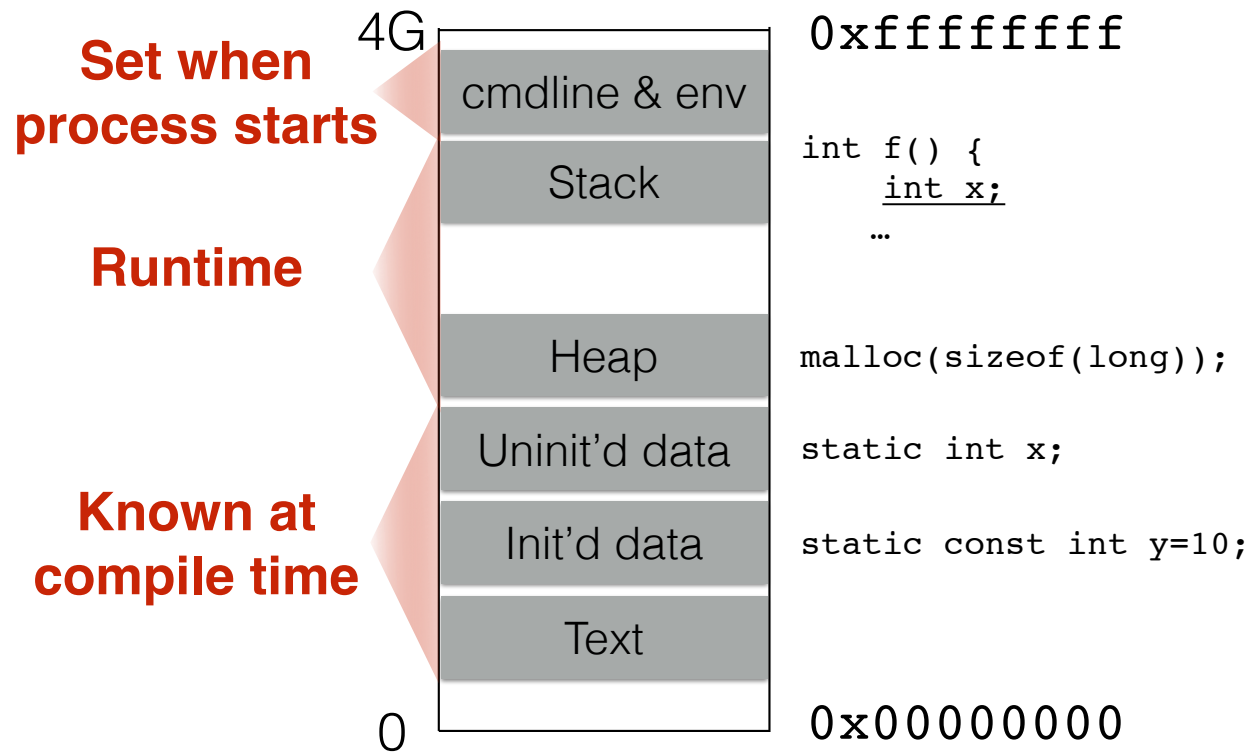
# All programs are stored in memory



The instructions themselves are in memory



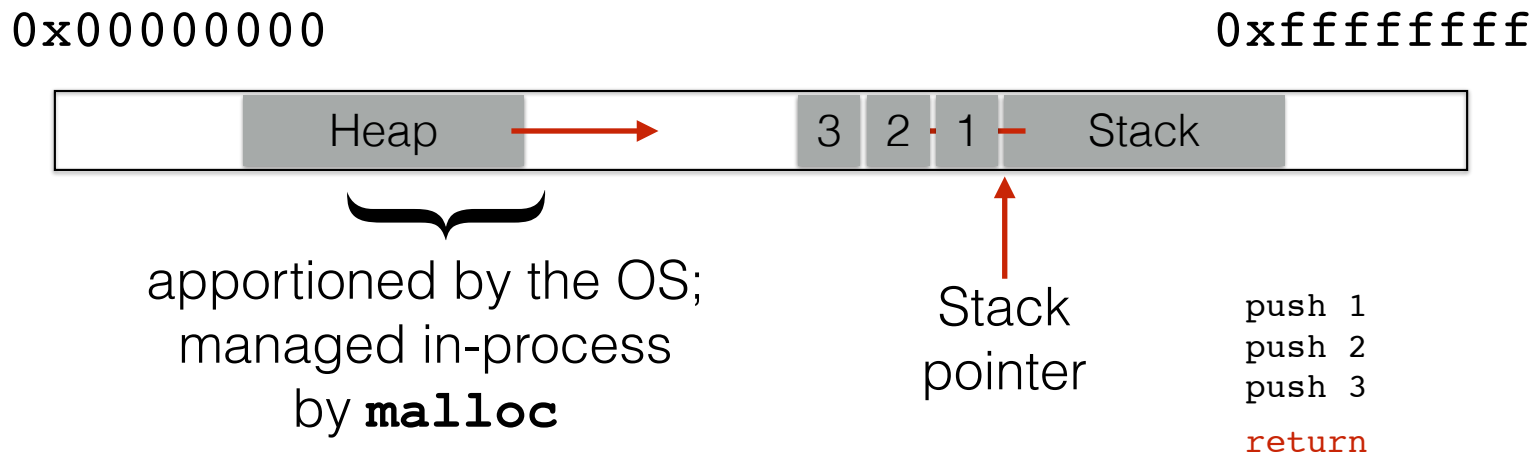
# Location of data areas



# Memory allocation

## Stack and heap grow in opposite directions

Compiler emits instructions  
adjust the size of the stack at run-time



**Focusing on the stack for now**

# Stack and function calls

- What happens when we **call** a function?
  - What data needs to be stored?
  - Where does it go?
- What happens when we **return** from a function?
  - What data needs to be *restored*?
  - Where does it come from?

# Basic stack layout

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4]
    int  loc2;
    ...
}
```

0xffffffff



**Local variables**  
pushed in the  
same order as  
they appear  
in the code

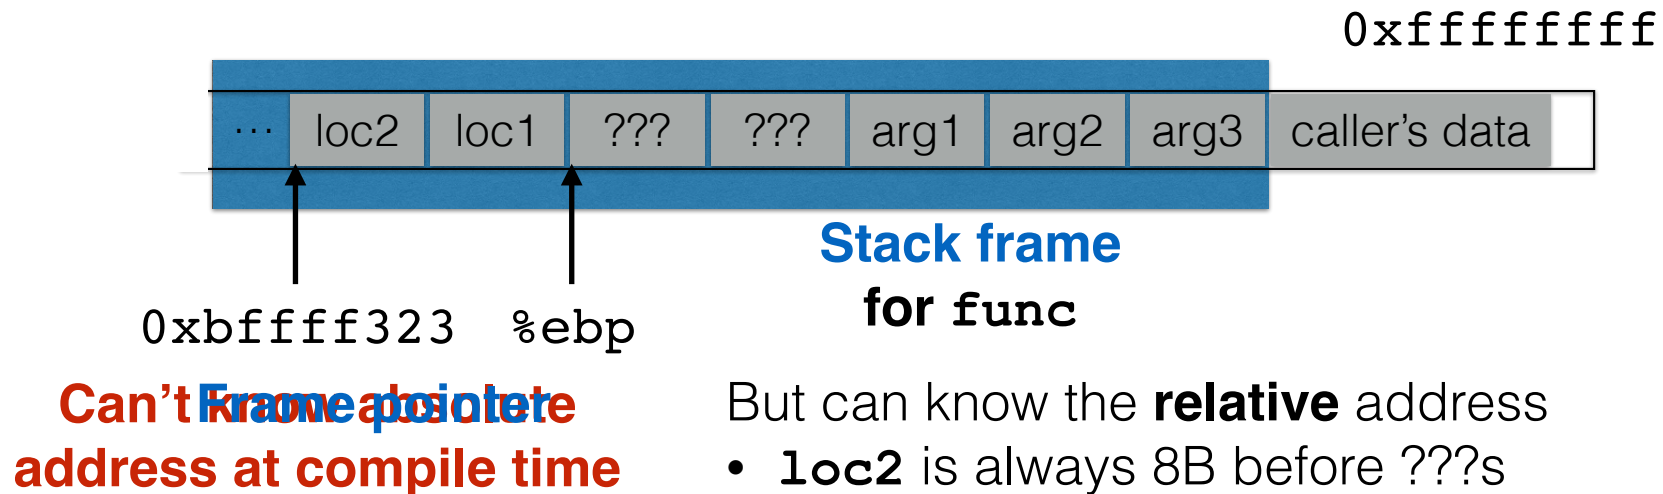
**Arguments**  
pushed in  
reverse order  
of code



# Accessing variables

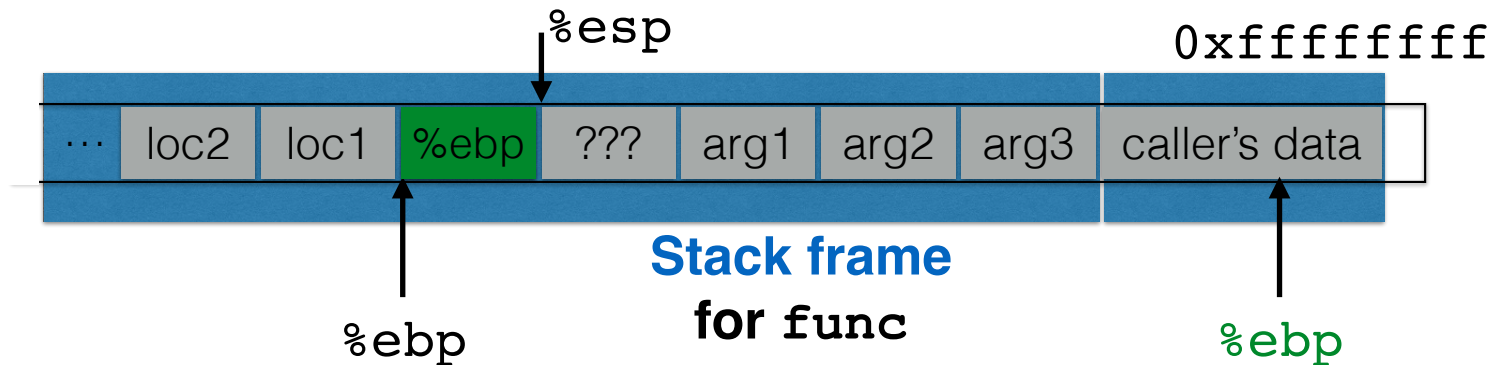
```
void func(char *arg1, int arg2, int arg3)
{
    ...
    loc2++;
    ...
}
```

**Q: Where is (this) loc2?**  
**A: -8(%ebp)**



# Returning from functions

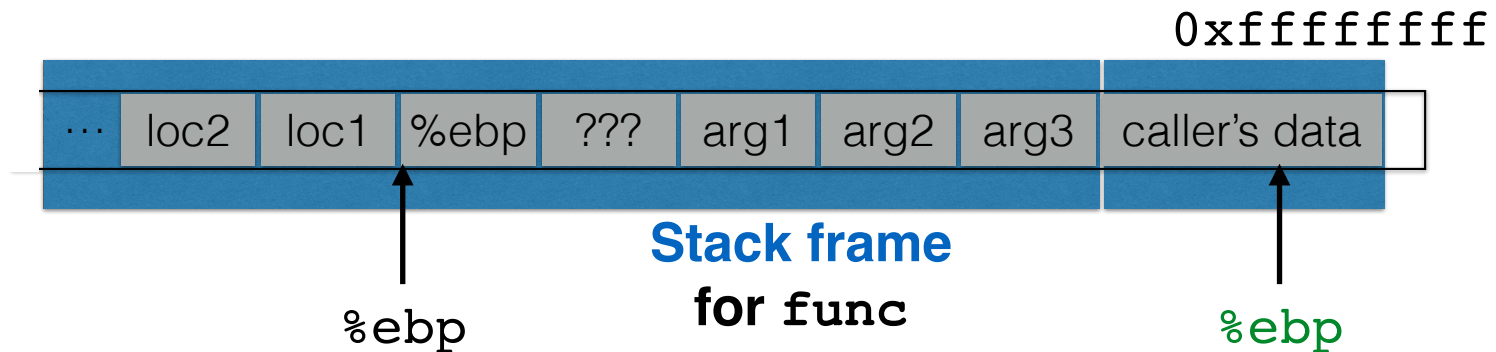
```
int main()  
{  
    ...  
    func("Hey", 10, -3);  
    ... Q: How do we restore %ebp?  
}
```



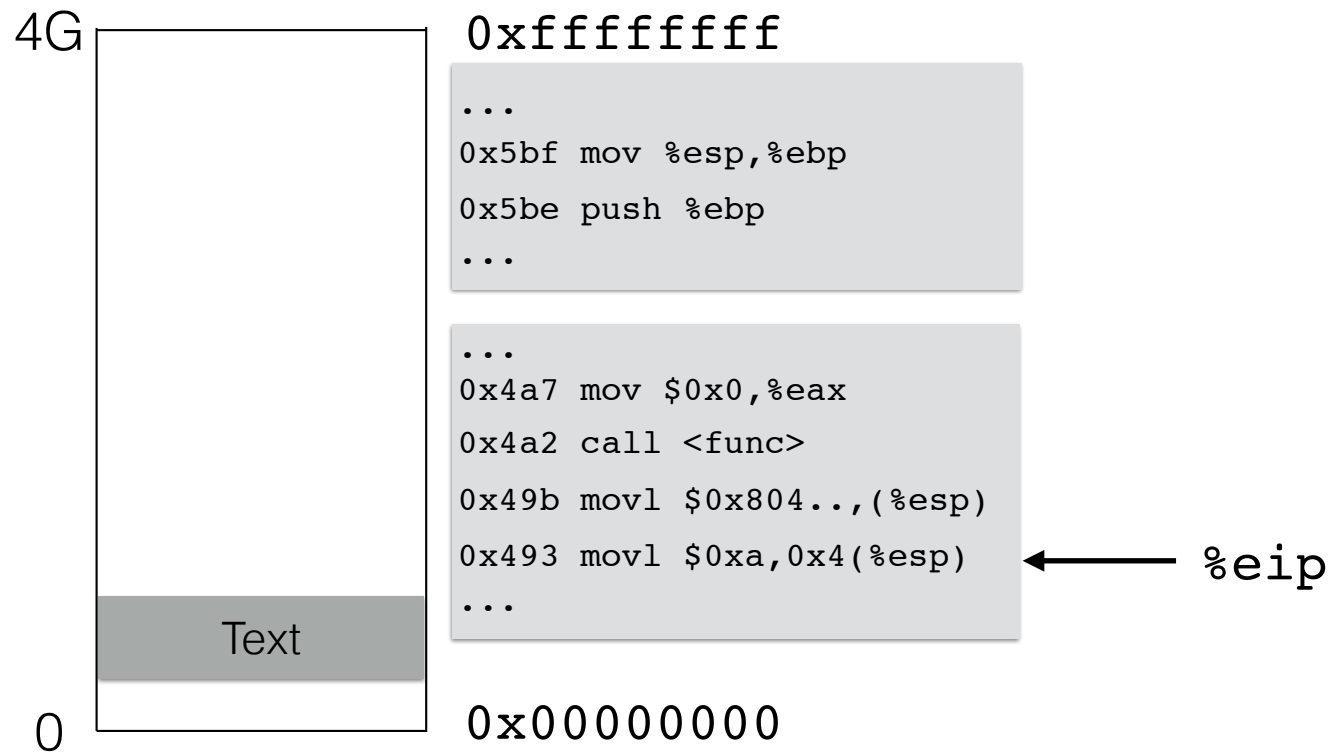
**Push %ebp before locals**  
**Set %ebp to current (%esp)**  
**Set %ebp to (%ebp) at return**

# Returning from functions

```
int main()  
{  
    ...  
    func("Hey", 10, -3);  
    ... Q: How do we resume here?  
}
```

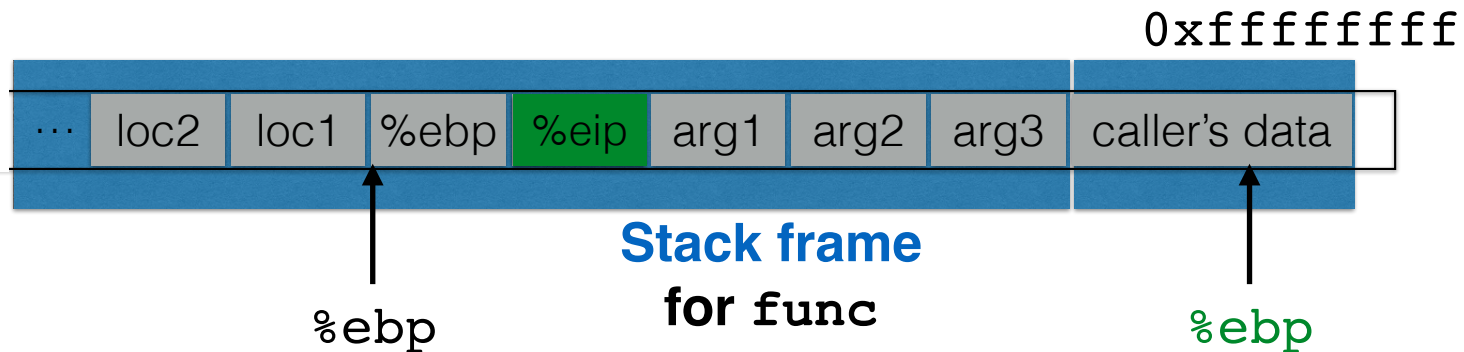


# Instructions in memory



# Returning from functions

```
int main()  
{  
    ...  
    func("Hey", 10, -3);  
    ... Q: How do we resume here?  
}
```



**Set `%eip` to 4(`%ebp`)  
at return**

**Push next `%eip`  
before call**

# Stack and functions: Summary

## Calling function:

1. **Push arguments** onto the stack (in reverse)
2. **Push the return address**, i.e., the address of the instruction you want run after control returns to you
3. **Jump to the function's address**

## Called function:

4. **Push the old frame pointer** onto the stack (%ebp)
5. **Set frame pointer** (%ebp) to where the end of the stack is right now (%esp)
6. **Push local variables** onto the stack

## Returning function:

7. **Reset the previous stack frame:** %ebp = (%ebp)
8. **Jump back to return address:** %eip = 4(%ebp)

Buffer overflows

# Buffer overflows from 10,000 ft

- **Buffer =**
  - Contiguous memory associated with a variable or field
  - Common in C
    - All strings are (NUL-terminated) arrays of `char`'s
- **Overflow =**
  - Put more into the buffer than it can hold
- **Where does the overflowing data go?**
  - Well, now that you are an expert in memory layouts...




# Benign outcome

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, str);
    ...
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

**Upon return, sets %ebp to 0x0021654d**



# Security-relevant outcome

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, str);
    if(authenticated) { ...
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

**Code still runs; user now 'authenticated'**

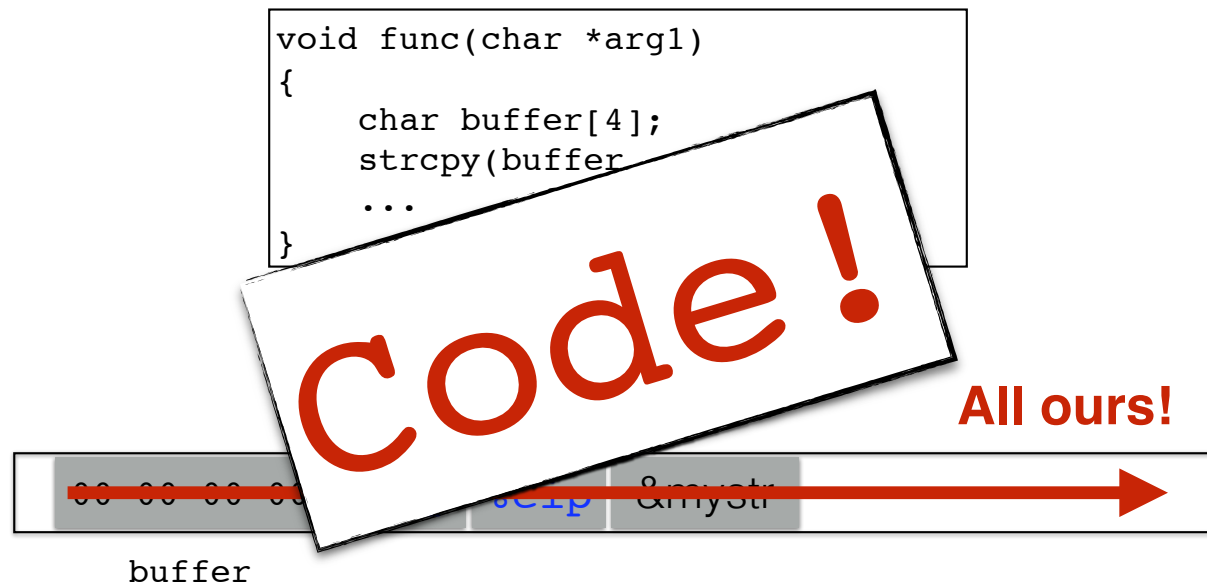
M e ! \0

	A u t h	4d 65 21 00	%ebp	%eip	&arg1	
--	---------	-------------	------	------	-------	--

buffer

authenticated

# Could it be worse?



**strcpy will let you write as much as you want (til a '\0')**  
**What could you write to memory to wreak havoc?**


# Aside: User-supplied strings

- These examples provide their own strings
- In reality **strings** come **from users** in myriad ways
  - **Text** input
  - **Packets**
  - **Environment variables**
  - **File** input...
- **Validating assumptions** about **user input** is extremely **important**
  - We will discuss it later, and throughout the course

# Code injection

# Code Injection: Main idea

```
void func(char *arg1)
{
    char buffer[4];
    sprintf(buffer, arg1);
    ...
}
```



**(1) Load my own code into memory**

**(2) Somehow get %eip to point to it**

# Challenge 1

## Loading code into memory

- It **must be the machine code** instructions (i.e., already compiled and ready to run)
- We have to be careful in how we construct it:
  - It **can't contain** any **all-zero bytes**
    - Otherwise, `sprintf` / `gets` / `scanf` / ... will stop copying
    - How could you write assembly to never contain a full zero byte?
  - It **can't use the loader** (we're injecting)
  - It **can't use the stack** (we're going to smash it)

# What code to run?

- Goal: **general-purpose shell**
  - Command-line prompt that gives attacker **general access to the system**
- The code to launch a shell is called **shellcode**



# Shellcode

```
#include <stdio.h>
int main( ) {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

**Assembly**

```
xorl %eax, %eax
pushl %eax
pushl $0x68732f2f
pushl $0x6e69622f
movl %esp, %ebx
pushl %eax
...
```

```
"\x31\xc0"
"\x50"
"\x68" "//sh"
"\x68" "/bin"
"\x89\xe3"
"\x50"
...
```

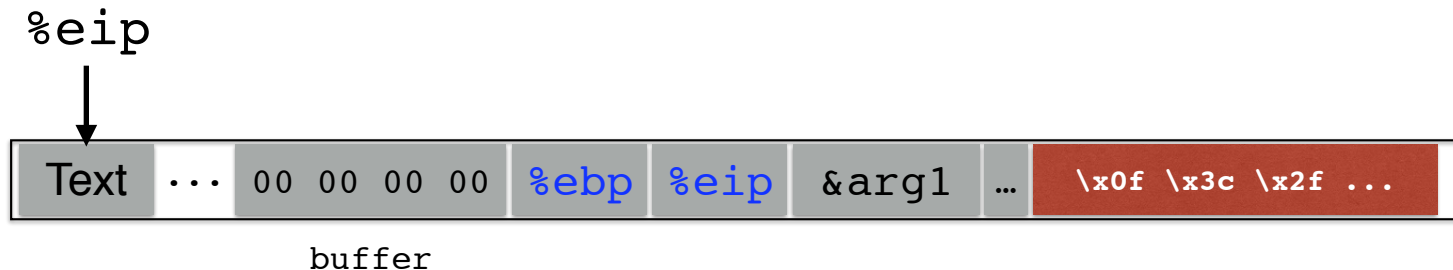
**Machine code**

(Part of)  
your  
input

# Challenge 2

## Getting injected code to run

- We can't insert a "jump into my code" instruction
- We don't know precisely where our code is



Recall

# Memory layout summary

## Calling function:

1. Push arguments onto the stack (in reverse)
2. Push the return address, i.e., the address of the instruction you want run after control returns to you
3. Jump to the function's address

## Called function:


4. Push the old frame pointer onto the stack (%ebp)
5. Set frame pointer (%ebp) to where the end of the stack is right now (%esp)
6. Push local variables onto the stack

## Returning function:

7. Reset the previous stack frame: %ebp = (%ebp)

8. **Jump back to return address:** %eip = 4(%ebp)

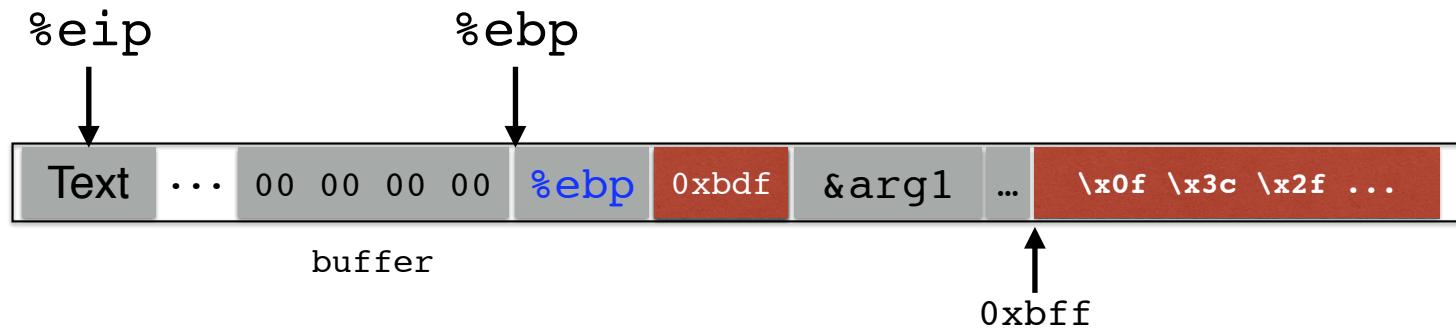
# Hijacking the saved %eip



**But how do we know the address?**

# Hijacking the saved %eip

**What if we are wrong?**



**This is most likely data,  
so the CPU will panic  
(Invalid Instruction)**

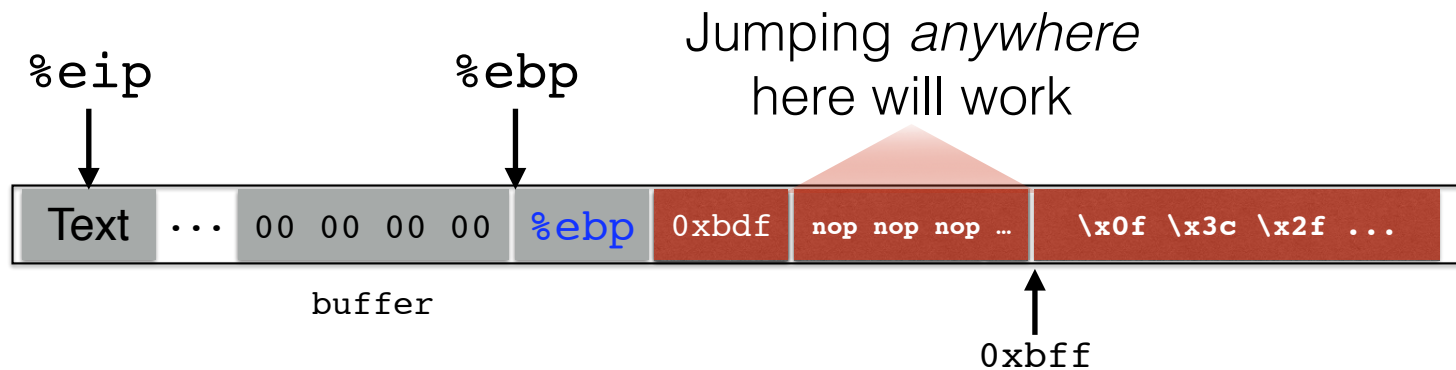
# Challenge 3

## Finding the return address

- If we don't have access to the code, we don't know how far the buffer is from the saved `%ebp`
- One approach: just try a lot of different values!
  - Worst case scenario: it's a 32 (or 64) bit memory space, which means  $2^{32}$  ( $2^{64}$ ) possible answers
- Without address randomization (discussed later):
  - The **stack always starts** from the same **fixed address**
  - The stack will grow, but usually it **doesn't grow very deeply** (unless the code is heavily recursive)

# Improving our chances: **nop** sleds

`nop` is a single-byte instruction  
(just moves to the next instruction)



**Now we improve our chances  
of guessing by a factor of #nops**

# Putting it all together

But it has to be *something*;  
we have to start writing wherever  
the input to `gets`/etc. begins.

