

Making secure software

How should we write software that will be secure?

- **Flawed approach:** Design and build software, and **ignore security at first**
 - Add security once the functional requirements are satisfied!
 - The software may have important vulnerability at the end
- **Better approach: Build security in** from the start
 - Incorporate security-minded thinking into all phases of the development process
 - Can avoid missing important security requirements

Development process

Many development processes; **four common** phases:

1. **Requirements!** Involves determining what the software should do, and not
2. **Design!** looks at how to structure the system to meet these requirements
3. **Implementation!** Involves actually writing code to implement the design
4. **Testing/assurance!** Involves checking that the implementation actually does what it's supposed to

Phases of development apply to the whole project, its individual components, and its refinements/iterations!

Where does **security engineering** fit in?

- **All phases!**

Security engineering

Phases

Requirements

Design!

Implementation!

Testing/assurance

Activities

Security Requirements

Abuse Cases

Architectural Risk Analysis

Security-oriented Design

Code Review (with tools)

Risk-based Security Tests

Penetration Testing

Note that different SD processes have different phases and artifacts, but all involve the basics above. We'll keep it simple and refer to these.

Running Example: On-line banking


What requirements might we have for this application?

- Account holders be able to deposit funds into their accounts
- Prevent other people who are not authorized by the account holder from withdrawing those funds
- The account holder accesses their funds whenever time they choose

Bob's



Alice's



Threat Modeling

The **threat model makes explicit the adversary's assumed powers!**

- Consequence: The threat model must match reality, otherwise the risk analysis of the system will be wrong

The threat model is **critically important**

- If you are not explicit about what the attacker can do, how can you assess whether your design will repel that attacker?
- This is part of **architectural risk analysis**


Example: Network User

An (anonymous) user that can connect to a service via the network

Can:

- **measure** the size and timing of requests and responses
- run **parallel sessions!**
- provide **malformed inputs, malformed messages**
- **drop or send extra messages**

- Example attacks: SQL injection, XSS, CSRF, buffer overrun/ROP payloads, ...



Example: Snooping User


Internet user **on the same network** as other users of some service

- For example, someone connected to an unencrypted Wi-Fi network at a coffee shop

Thus, can additionally

- **Read/measure** others' messages,
- **Intercept, duplicate, and modify** messages!

Example attacks: Session hijacking (and other data theft), **privacy-violating side-channel attack, denial of service**



Example: Co-located User


Internet user **on the same machine** as other users of some service

- E.g., malware installed on a user's laptop

Thus, can additionally

- **Read/write** user's **files** (e.g., cookies) and memory
- Snoop keypresses and other events
- Read/write the user's **display** (e.g., to spoof)

Example attacks: Password theft (and other credentials/secrets)



Threat-driven Design

Different threat models will elicit different responses

Network-only attackers implies **message traffic is safe**

- No need to encrypt communications
- This is what *telnet* remote login software assumed

Snooping attackers means **message traffic is visible**

- So use encrypted wifi (link layer), encrypted network layer (IPsec), or encrypted application layer (SSL)
 - Which is most appropriate for your system?

Co-located attacker can **access local files, memory**

- Cannot store unencrypted secrets, like passwords

Bad Model = Bad Security

Any **assumptions** you make in your model are potential **holes that the adversary can exploit!**

E.g.: **Assuming no snooping users no longer valid**

- Prevalence of wi-fi networks in most deployments

Other mistaken assumptions

- **Assumption: Encrypted traffic carries no information**

- Not true! By analyzing the size and distribution of messages, you can infer application state

- **Assumption: Timing channels carry little information**

- Not true! Timing measurements of previous RSA implementations could be used eventually reveal a remote SSL secret key

Finding a good model

Compare against similar systems

- What attacks does their design contend with?

Understand past attacks and attack patterns

- How do they apply to your system?

Challenge assumptions in your design

- What happens if an assumption is untrue?
 - What would a breach potentially cost you?
- How hard would it be to get rid of an assumption, allowing for a stronger adversary?
 - What would that development cost?

Security Requirements

Security Requirements

Software requirements typically about **what the software should do**

Security requirements **Security-related** **goals** (or **policies**)

- **Example:** One user's bank account balance should not be learned by, or modified by, another user, unless authorized

Required **mechanisms** **for enforcing them**

- **Example:**
 - Users identify themselves using passwords,
 - Passwords must be “strong,” and
 - The password database is only accessible to login program.

Typical *Kinds* of Requirements

Policies

- **Confidentiality** (and Privacy and Anonymity)
- **Integrity**
- **Availability**

Supporting mechanisms

- **Authentication**
- **Authorization**
- **Auditability**

Privacy and Confidentiality

Definition: **Sensitive information not leaked** to unauthorized parties

- Called *privacy* for individuals, *confidentiality* for data
- **Example** policy: bank account status (including balance) known only to the account owner

Leaking **directly** or via **side channels!**

- **Example:** manipulating the system to directly display Bob's bank balance to Alice
- **Example:** determining Bob has an account at Bank A according to shorter delay on login failure

Secrecy vs. Privacy?

Anonymity

A specific **kind of privacy**

Example: Non-account holders should be able to browse the bank informational site without being tracked

- Here *the adversary is the bank*
- The previous examples considered other account holders as possible adversaries

Integrity

Definition: **Sensitive information not damaged** by (computations acting on behalf of) unauthorized parties

Example: Only the account owner can authorize withdrawals from her account

Violations of integrity can also be **direct** or **indirect**

- **Example:** Being able specifically withdraw from the account vs. confusing the system into doing it
 - For example, by using a cross-site request forgery (CSRF).

Availability

Definition: A system is **responsive to requests**

Example: a user may always access her account for balance queries or withdrawals

Denial of Service (DoS) attacks attempt to **compromise availability** by busying a system with useless work or cutting off network access

Supporting mechanisms

Leslie Lamport's defines **gold standard** mechanisms provided by a system to enforce its requirements

- **Authentication**
- **Authorization**
- **Audit**

The gold standard is **both requirement and design**

- The *sorts of policies* that are authorized *determines* the *authorization mechanism*
- The *sorts of users* a system has *determines* how they should be *authenticated*

Authentication

What is the **subject of security policies**?

- Need to define a ***notion of identity*** and a way to ***connect an action with an identity!***
 - By who we mean a **principal**. That is, it could also be a human being. Or it could be some service or a computer program.

How can system tell a user is who he says he is?

- What (only) he **knows** (e.g., password)
- What he **is** (e.g., biometric)
- What he **has** (e.g., smartphone)
- Authentication mechanisms that employ more than one of these factors are called **multi-factor authentication**
 - E.g., bank may employ passwords and text of a special code to a user's smart phone

Authorization

Defines **when a principal may perform an action!**

Example: Bob is authorized to access his own account, but not Alice's account

There are a wide variety of **policies** that define what actions might be authorized

- E.g., access control policies, which could be
 - originator based: the users performing particular actions , like Alice or Bob.
 - role-based: a bank teller or a bank manager may be allowed to perform certain things.
 - user-based: Bob could originate a policy that allows Alice to withdraw a fixed amount of money from his account.
 - etc.

Audit

Retain enough information to be able to **determine the circumstances of a breach or misbehavior** (*or establish one did not occur*)

Such information, often stored in **log files**, must be **protected from tampering**, and from access that might violate other policies

Example: Every account-related action is logged locally and mirrored at a separate site

Defining Security Requirements

Many processes for deciding security requirements

Example: **General policy concerns**

- Due to **regulations**/standards
 - E.g. HIPAA: Health Insurance Portability and Accountability Act in US
- Due **organizational values** (e.g., valuing privacy)

Example: **Policy arising from threat modeling**

- Which **attacks** cause the **greatest concern**?
 - Who are the likely adversaries and what are their goals and methods?
- Which **attacks** have **already occurred**?
 - Within the organization, or elsewhere on related systems?

Abuse Cases

Abuse cases illustrate security requirements

Where use cases describe what a system *should* do, **abuse cases describe what it should *not* do!**

Example **use case**: The system allows bank managers to modify an account's interest rate

Example **abuse case**: A user is able to spoof being a manager and thereby change the interest rate on an account

Defining Abuse Cases

Using attack patterns and likely scenarios, construct cases in which an **adversary's exercise of power** could **violate a security requirement**

- Based on the threat model
- What might occur if a security measure was removed?

Example: *Co-located attacker* steals password file and learns all user passwords

- Possible if password file is not encrypted

Example: *Snooping attacker* replays a captured message, effecting a bank withdrawal

- Possible if messages are have no *nonce*